

AI505
Optimization

First-Order Methods

Marco Chiarandini

Department of Mathematics & Computer Science
University of Southern Denmark

Outline

1. Gradient Descent
2. Conjugate Descent
3. Accelerated Descents

Descent Direction Methods

How to select the descent direction?

- first-order methods that rely on gradient
- second-order methods that rely on Hessian information

Advantages of first order methods:

- cheap iterations: good for small and large scale optimization
- helpful because easy to warm restart

Limitations of first order methods:

- not hard to find challenging instances for them.
- can converge slowly.

Outline

1. Gradient Descent
2. Conjugate Descent
3. Accelerated Descents

Gradient Descent

The **steepest descent** direction at \mathbf{x}_k , at k th iteration of a **local descent iterative method**, is the one opposite to the gradient (**gradient descent**):

$$\mathbf{d}_k = -\frac{\nabla f(\mathbf{x}_k)}{\|\nabla f(\mathbf{x}_k)\|}$$

Guaranteed to lead to improvement if:

- f is smooth
- step size is sufficiently small
- \mathbf{x}_k is not a stationary point (ie, $\nabla f(\mathbf{x}_k) = 0$)

Gradient Descent: Example

- Suppose we have

$$f(\mathbf{x}) = x_1 x_2$$

- The gradient is $\nabla f = [x_2^2, 2x_1 x_2]$
- $\mathbf{x}_k = [1, 2]$

$$\mathbf{d}_{k+1} = -\frac{\nabla f(\mathbf{x}_k)}{\|\nabla f(\mathbf{x}_k)\|} = \frac{[-4, -4]}{\sqrt{16 + 16}} = \left[-\frac{1}{\sqrt{2}}, -\frac{1}{\sqrt{2}} \right]$$

Implementation

Gradient Descent
Conjugate Descent
Accelerated Descents

```
class DescentMethod:  
    alpha: float  
  
class GradientDescent(DescentMethod):  
    def __init__(self, f, grad, x, alpha):  
        self.alpha = alpha  
  
    def step(self, f, grad, x):  
        alpha, g = self.alpha, grad(x)  
        return x - alpha * g
```

Gradient Descent

Theorem: The next direction is orthogonal to the current direction.

Proof:

$$\alpha_k^* = \underset{\alpha}{\operatorname{argmin}} f(\mathbf{x}_k + \alpha \mathbf{d}_k)$$

$$\nabla f(\mathbf{x}_k + \alpha_k^* \mathbf{d}_k) = \nabla_{\mathbf{d}_k} f(\mathbf{x}_k + \alpha_k^* \mathbf{d}_k) = 0$$

because α_k^* is minimum

$$\nabla f(\mathbf{x}_k + \alpha_k^* \mathbf{d}_k)^T \mathbf{d}_k = 0$$

because directional derivative: $\nabla_{\mathbf{s}} f(\mathbf{x}) = \nabla f(\mathbf{x})^T \mathbf{s}$

$$\mathbf{d}_{k+1} = -\frac{\nabla f(\mathbf{x}_k + \alpha_k^* \mathbf{d}_k)}{\|\nabla f(\mathbf{x}_k + \alpha_k^* \mathbf{d}_k)\|}$$

gradient descent

$$\mathbf{d}_{k+1} \cdot \mathbf{d}_k = -\frac{\nabla f(\mathbf{x}_k + \alpha_k^* \mathbf{d}_k)}{\|\nabla f(\mathbf{x}_k + \alpha_k^* \mathbf{d}_k)\|} \cdot \mathbf{d}_k = 0$$

$$\mathbf{d}_{k+1}^T \mathbf{d}_k = 0 \implies \mathbf{d}_{k+1} \perp \mathbf{d}_k$$

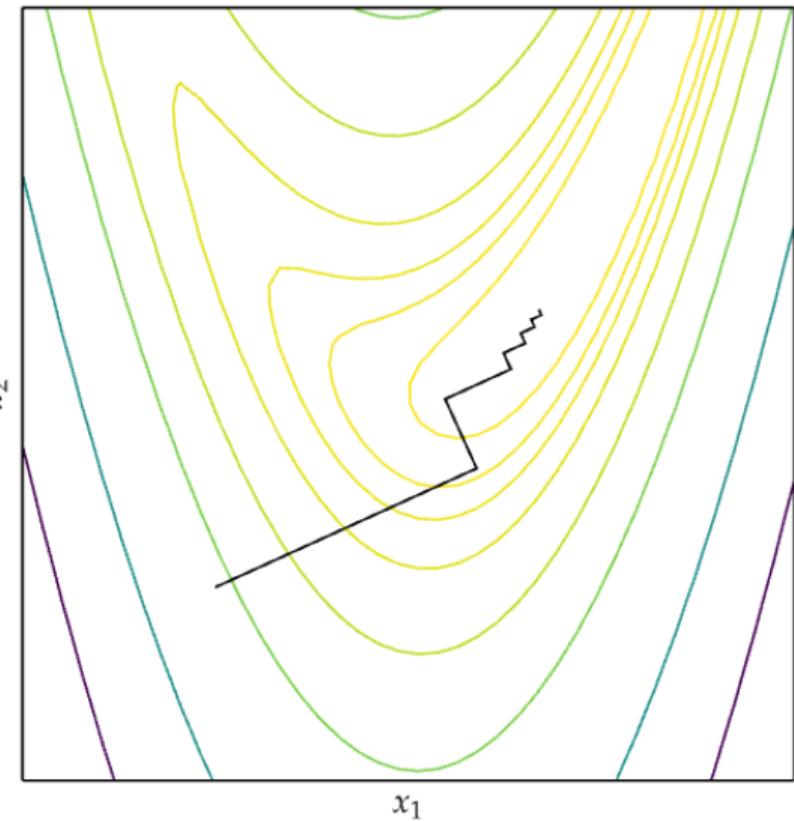


Gradient Descent: Example

2D Rosenbrock function

$$f(x, y) = (a - x)^2 + b(y - x^2)^2$$

Narrow valleys not aligned with gradient can be a problem



Outline

1. Gradient Descent
2. Conjugate Descent
3. Accelerated Descents

Conjugate Gradient

[Hestenes and Stiefel, 1950s]

For A symmetric positive definite:

$$Ax = b \iff \underset{x}{\text{minimize}} f(x) \stackrel{\text{def}}{=} \frac{1}{2}x^T Ax - b^T x$$

$$\nabla f(x) = Ax - b \stackrel{\text{def}}{=} r(x)$$

Conjugate Direction

Def.: A set of nonzero vectors $\{\mathbf{d}_0, \mathbf{d}_1, \dots, \mathbf{d}_\ell\}$ is said to be **conjugate** with respect to the symmetric positive definite matrix A if

$$\mathbf{d}_i^T A \mathbf{d}_j = 0, \quad \text{for all } i \neq j$$

(the vectors are linearly independent. Generally, not orthogonal.)

Theorem: Given an arbitrary $\mathbf{x}_0 \in \mathbb{R}^n$ and a set of conjugate vectors $\{\mathbf{d}_0, \mathbf{d}_1, \dots, \mathbf{d}_{n-1}\}$ the sequence $\{\mathbf{x}_k\}$ generated by

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{d}_k$$

where α_k is the analytical solution of $\min_{\alpha} f(\mathbf{x}_k + \alpha \mathbf{d}_k)$ given by:

$$\alpha_k = -\frac{\mathbf{r}_k^T \mathbf{d}_k}{\mathbf{d}_k^T A \mathbf{d}_k}$$

(aka, **conjugate direction algorithm**) converges to the solution \mathbf{x}^* of the linear system and minimization problem in at most n steps.

Proof:

$$\min_{\alpha} f(\mathbf{x}_k + \alpha \mathbf{d}_k)$$

We can compute the derivative with respect to α :

$$\begin{aligned}\frac{\partial}{\partial \alpha} f(\mathbf{x} + \alpha \mathbf{d}) &= \frac{\partial}{\partial \alpha} (\mathbf{x} + \alpha \mathbf{d})^T A (\mathbf{x} + \alpha \mathbf{d}) - \mathbf{b}^T (\mathbf{x} + \alpha \mathbf{d}) (+c) \\ &= \mathbf{d}^T A (\mathbf{x} + \alpha \mathbf{d}) - \mathbf{d}^T \mathbf{b} \\ &= \mathbf{d}^T (A \mathbf{x} - \mathbf{b}) + \alpha \mathbf{d}^T A \mathbf{d}\end{aligned}$$

Setting $\frac{\partial f(\mathbf{x} + \alpha \mathbf{d})}{\partial \alpha} = 0$ results in:

$$\alpha_k = -\frac{\mathbf{d}_k^T (A \mathbf{x}_k - \mathbf{b})}{\mathbf{d}_k^T A \mathbf{d}_k} = -\frac{\mathbf{d}_k^T r(\mathbf{x}_k)}{\mathbf{d}_k^T A \mathbf{d}_k} \quad (1)$$

- Since the directions $\{\mathbf{d}_k\}$ are linearly independent, they must span the whole space \mathbb{R}^n . Hence, there is a set of scalars σ_k such that:

$$\mathbf{x}^* - \mathbf{x}_0 = \sigma_0 \mathbf{d}_0 + \sigma_1 \mathbf{d}_1 + \dots + \sigma_{n-1} \mathbf{d}_{n-1}$$

- By premultiplying this expression by $\mathbf{d}_k^T A$ and using the conjugacy property, we obtain:

$$\sigma_k = \frac{\mathbf{d}_k^T A(\mathbf{x}^* - \mathbf{x}_0)}{\mathbf{d}_k^T A \mathbf{d}_k} \quad (2)$$

- If \mathbf{x}_k is generated by conjugate direction algorithm, then we have

$$\mathbf{x}_k = \mathbf{x}_0 + \alpha_0 \mathbf{d}_0 + \alpha_1 \mathbf{d}_1 + \dots + \alpha_k \mathbf{d}_{k-1}$$

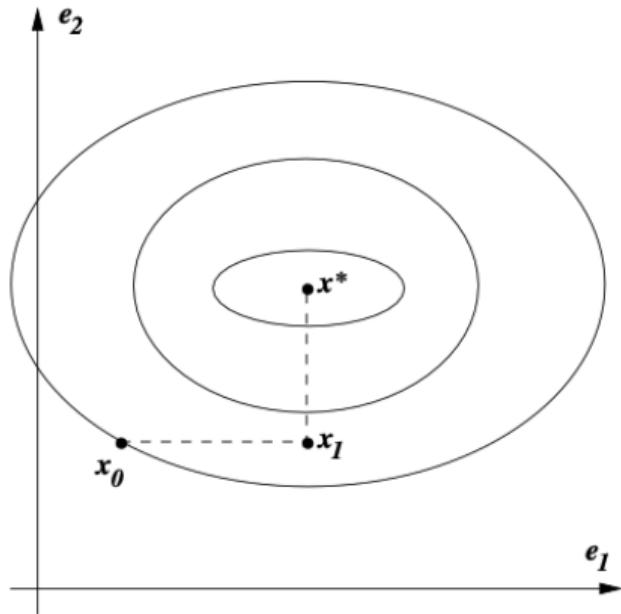
- By premultiplying this expression by $\mathbf{d}_k^T A$ and using the conjugacy property, we have that

$$\mathbf{d}_k^T A(\mathbf{x}_k - \mathbf{x}_0) = 0$$

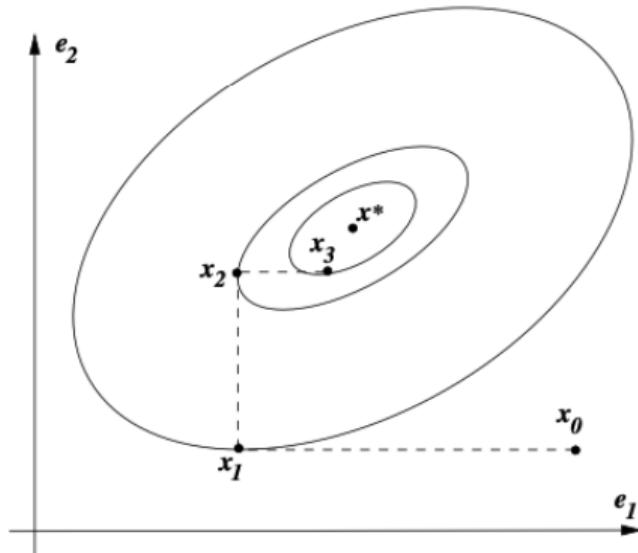
- and therefore

$$\begin{aligned} \mathbf{d}_k^T A(\mathbf{x}^* - \mathbf{x}_0) &= \mathbf{d}_k^T A(\mathbf{x}^* - \mathbf{x}_k + \mathbf{x}_k - \mathbf{x}_0) = \mathbf{d}_k^T A(\mathbf{x}^* - \mathbf{x}_k) + \mathbf{d}_k^T A(\mathbf{x}_k - \mathbf{x}_0) = \\ &= \mathbf{d}_k^T A(\mathbf{x}^* - \mathbf{x}_k) = \mathbf{d}_k^T (\mathbf{b} - A\mathbf{x}_k) = -\mathbf{d}_k^T \mathbf{r}_k. \end{aligned}$$

- Using this result in (2) and comparing with (1) we conclude $\alpha_k = \sigma_k$. □



If the matrix A is diagonal, the contours of the function $f(\cdot)$ are ellipses whose axes are aligned with the coordinate directions



If A is not diagonal, its contours are elliptical, but they are usually not aligned with the coordinate directions.

Transform the problem to make A diagonal and minimize along the coordinate directions.

Conjugate Gradient Method

- The **conjugate gradient method** is a **conjugate direction method** with the property: In generating its set of conjugate vectors, it can compute a new vector \mathbf{d}_k by using only the previous vector \mathbf{d}_{k-1} . Hence, little storage and computation requirements.

$$\mathbf{d}_k = -\mathbf{r}_k + \beta_k \mathbf{d}_{k-1}$$

where β_k is to be determined such that \mathbf{d}_{k-1} and \mathbf{d}_k must be conjugate with respect to A . By premultiplying by $\mathbf{d}_{k-1}^T A$ and imposing that $\mathbf{d}_{k-1}^T A \mathbf{d}_k = 0$ we find that

$$\beta_k = \frac{\mathbf{r}_k^T A \mathbf{d}_{k-1}}{\mathbf{d}_{k-1}^T A \mathbf{d}_{k-1}}$$

- Larger values of β indicate that the previous descent direction contributes more strongly.
- \mathbf{d}_0 is commonly chosen to be the steepest descent direction at \mathbf{x}_0
- Advantage with respect to steepest descent: implicitly reuses previous information about the function and thus better convergence.

Algorithm CG

Basic version:

Input: f, \mathbf{x}_0

Output: \mathbf{x}^*

Set $\mathbf{r}_0 \leftarrow \mathbf{A}\mathbf{x}_0 - \mathbf{b}$, $\mathbf{d}_0 \leftarrow \mathbf{r}_0$, $k \leftarrow 0$;

while $\mathbf{r}_k \neq 0$ **do**

$$\alpha_k \leftarrow -\frac{\mathbf{d}_k^T \mathbf{r}(\mathbf{x}_k)}{\mathbf{d}_k^T \mathbf{A}\mathbf{d}_k};$$

$$\mathbf{x}_{k+1} \leftarrow \mathbf{x}_k + \alpha_k \mathbf{d}_k;$$

$$\mathbf{r}_{k+1} \leftarrow \mathbf{A}\mathbf{x}_{k+1} - \mathbf{b};$$

$$\beta_{k+1} \leftarrow \frac{\mathbf{r}_{k+1}^T \mathbf{A}\mathbf{d}_k}{\mathbf{d}_k^T \mathbf{A}\mathbf{d}_k};$$

$$\mathbf{d}_{k+1} \leftarrow -\mathbf{r}_{k+1} + \beta_{k+1} \mathbf{d}_k;$$

$$k \leftarrow k + 1;$$

Computationally improved version:

Input: f, \mathbf{x}_0

Output: \mathbf{x}^*

Set $\mathbf{r}_0 \leftarrow \mathbf{A}\mathbf{x}_0 - \mathbf{b}$, $\mathbf{d}_0 \leftarrow \mathbf{r}_0$, $k \leftarrow 0$;

while $\mathbf{r}_k \neq 0$ **do**

$$\alpha_k \leftarrow -\frac{\mathbf{r}(\mathbf{x}_k)^T \mathbf{r}(\mathbf{x}_k)}{\mathbf{d}_k^T \mathbf{A}\mathbf{d}_k};$$

$$\mathbf{x}_{k+1} \leftarrow \mathbf{x}_k + \alpha_k \mathbf{d}_k;$$

$$\mathbf{r}_{k+1} \leftarrow \mathbf{r}_k + \alpha_k \mathbf{A}\mathbf{d}_k;$$

$$\beta_{k+1} \leftarrow \frac{\mathbf{r}_{k+1}^T \mathbf{r}_{k+1}}{\mathbf{r}_k^T \mathbf{r}_k};$$

$$\mathbf{d}_{k+1} \leftarrow -\mathbf{r}_{k+1} + \beta_{k+1} \mathbf{d}_k;$$

$$k \leftarrow k + 1;$$

- we never need to know the vectors \mathbf{x} , \mathbf{r} , and \mathbf{d} for more than the last two iterations.
- major computational tasks: the matrix–vector product $\mathbf{A}\mathbf{d}_k$, inner products $\mathbf{d}_k^T \mathbf{A}\mathbf{d}_k$ and $\mathbf{r}_{k+1}^T \mathbf{r}_{k+1}$, and three vector sums

NonLinear Conjugate Gradient Methods

- The conjugate gradient method can be applied to nonquadratic functions as well.
- Smooth, continuous functions behave like quadratic functions close to a local minimum
- but! we do not know the value of A that best approximates f around x_k . Instead, several choices for β_k tend to work well:
- Two changes:
 - α_k is computed by solving an approximate line search
 - the residual r , (it was simply the gradient of f), must be replaced by the gradient of the nonlinear objective f .

NonLinear Conjugate Gradient Methods

Fletcher-Reeves Method:

Input: f, \mathbf{x}_0

Output: \mathbf{x}^*

Evaluate $f_0 = f(\mathbf{x}_0), \nabla f_0 = \nabla f(\mathbf{x}_0)$;

Set $\mathbf{d}_0 \leftarrow -\nabla f_0, k \leftarrow 0$;

while $\nabla f_k \neq 0$ **do**

 Compute α_k by line search and set

$\mathbf{x}_{k+1} \leftarrow \mathbf{x}_k + \alpha_k \mathbf{d}_k$;

 Evaluate ∇f_{k+1} ;

$\beta_{k+1}^{FR} \leftarrow \frac{\nabla f_{k+1}^T \nabla f_{k+1}}{\nabla f_k^T \nabla f_k}$;

$\mathbf{d}_{k+1} \leftarrow -\nabla f_{k+1} + \beta_{k+1}^{FR} \mathbf{d}_k$;

$k \leftarrow k + 1$;

Polak-Ribière:

Input: f, \mathbf{x}_0

Output: \mathbf{x}^*

Evaluate $f_0 = f(\mathbf{x}_0), \nabla f_0 = \nabla f(\mathbf{x}_0)$;

Set $\mathbf{d}_0 \leftarrow -\nabla f_0, k \leftarrow 0$;

while $\nabla f_k \neq 0$ **do**

 Compute α_k by line search and set

$\mathbf{x}_{k+1} \leftarrow \mathbf{x}_k + \alpha_k \mathbf{d}_k$;

 Evaluate ∇f_{k+1} ;

$\beta_{k+1}^{PR} \leftarrow \frac{\nabla f_{k+1}^T (\nabla f_{k+1} - \nabla f_k)}{\nabla f_k^T \nabla f_k}$;

$\mathbf{d}_{k+1} \leftarrow -\nabla f_{k+1} + \beta_{k+1}^{FR} \mathbf{d}_k$;

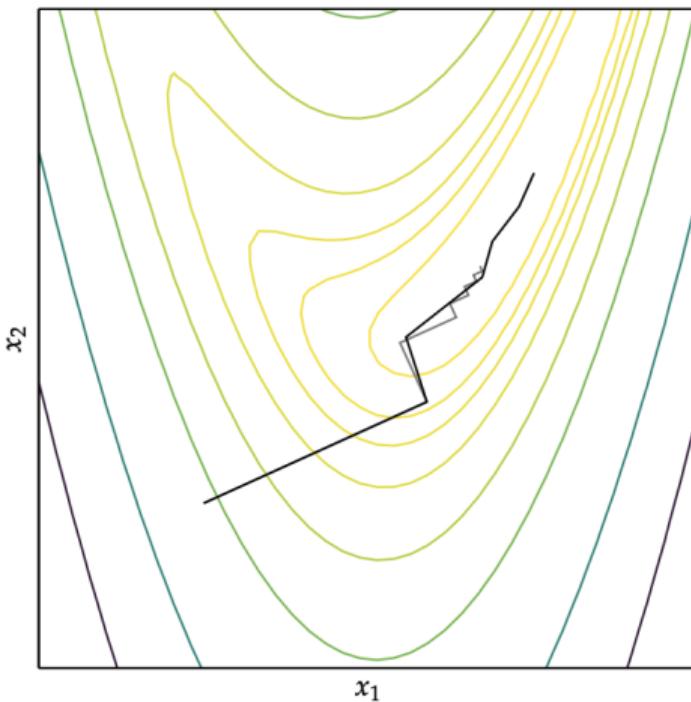
$k \leftarrow k + 1$;

PR with:

$$\beta_{k+1}^+ = \max\{\beta_{k+1}^{PR}, 0\}$$

becomes PR⁺ and guaranteed to converge (satisfies first Wolfe conditions).

The conjugate gradient method with the Polak-Ribi re update. Gradient descent is shown in gray.

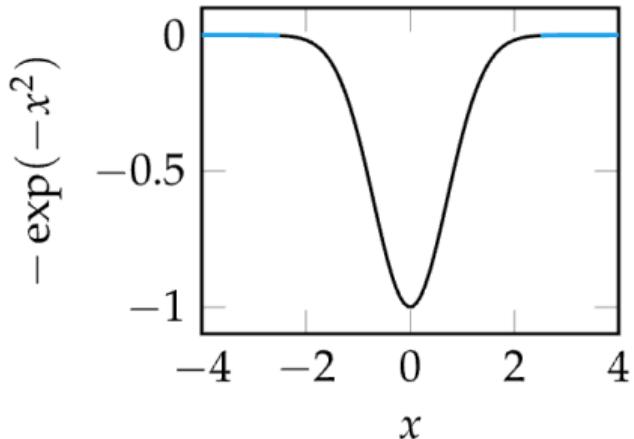


Outline

1. Gradient Descent
2. Conjugate Descent
3. Accelerated Descents

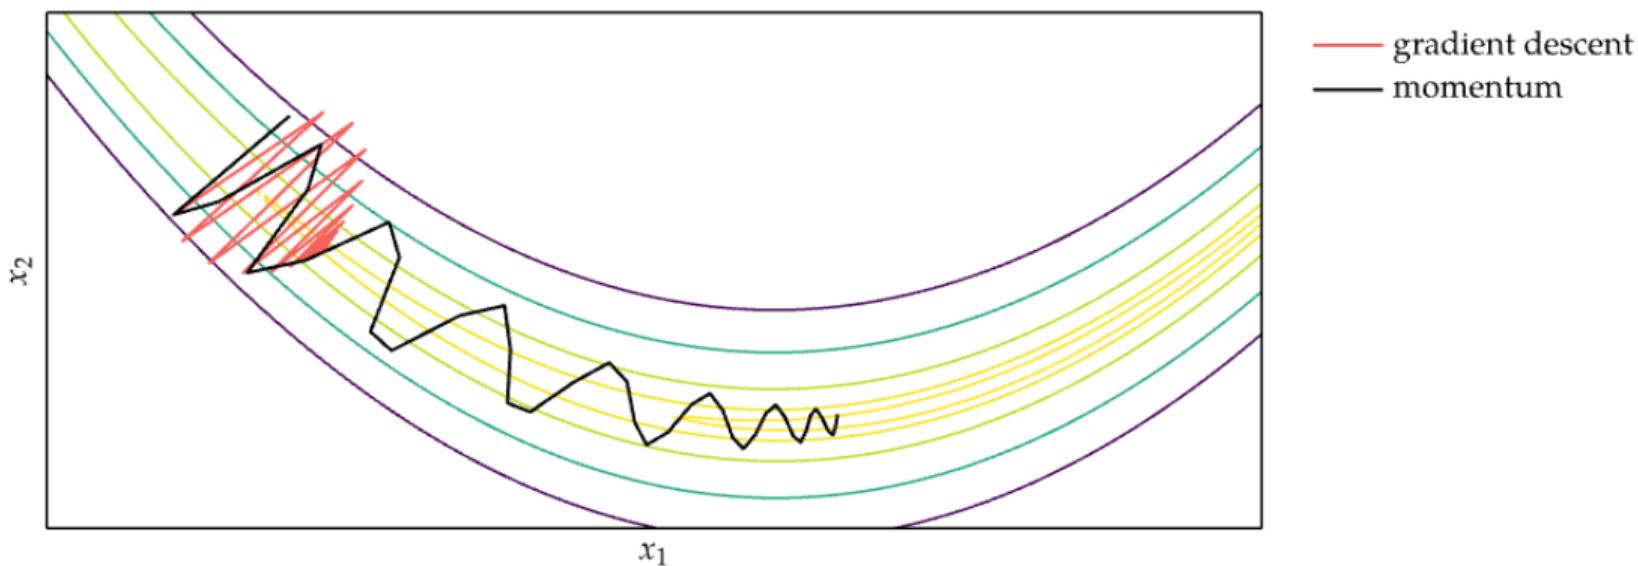
Accelerated Descents

- Addresses common convergence issues
- Some functions have regions with very small gradients (flat surface) where gradient descent gets stuck



Momentum

Rosenbrock function with $b = 100$



Momentum overcomes these issues by replicating the effect of physical momentum

Momentum

Momentum update equations:

$$\begin{aligned}\mathbf{v}_{k+1} &= \beta \mathbf{v}_k - \alpha \nabla f(\mathbf{x}_k) \\ \mathbf{x}_{k+1} &= \mathbf{x}_k + \mathbf{v}_{k+1}\end{aligned}$$

```
import numpy as np

class Momentum(DescentMethod):
    alpha: float # learning rate
    beta: float # momentum decay
    v: np.array # momentum

    def __init__(self, alpha, beta, f, grad, x):
        self.alpha = alpha
        self.beta = beta
        self.v = np.zeros_like(x)

    def step(self, grad, x):
        self.v = self.beta * self.v - self.alpha * ↵
            ↵grad(x)
        return x + self.v
```

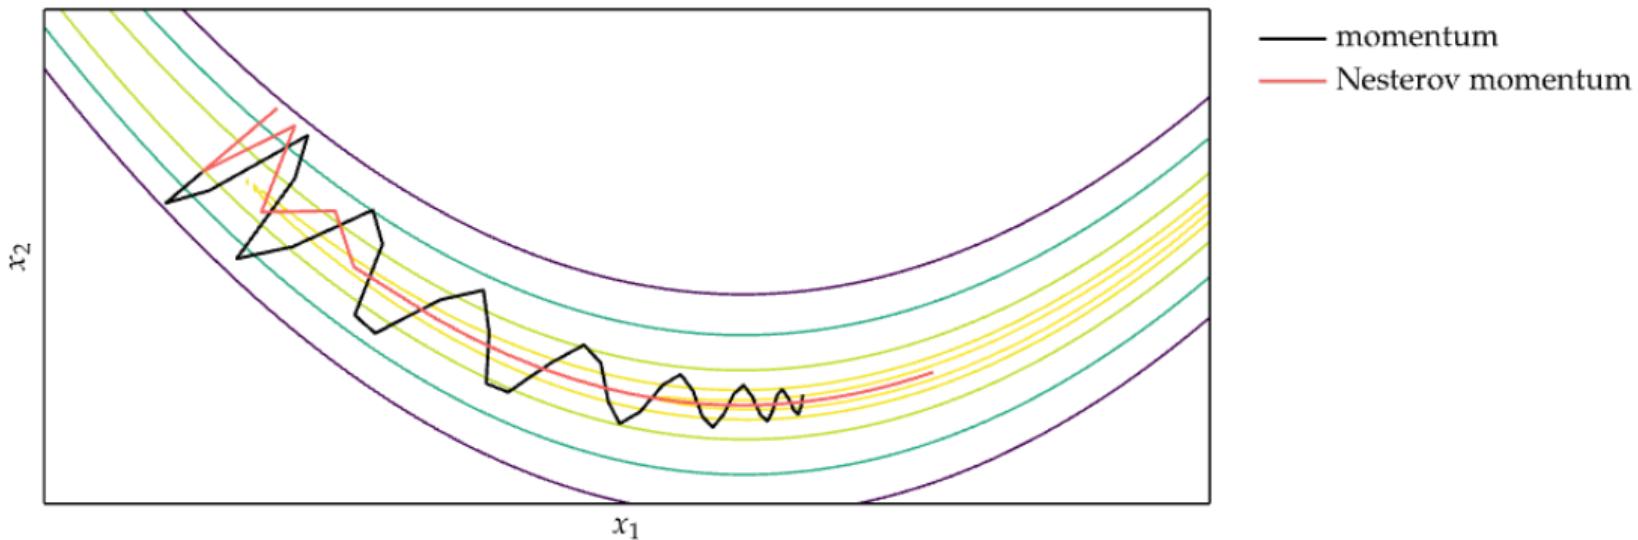
Nesterov Momentum

Issue of momentum: steps do not slow down enough at the bottom of a valley, overshoot.

Nesterov Momentum update equations:

$$\mathbf{v}_{k+1} = \beta \mathbf{v}_k - \alpha \nabla f(\mathbf{x}_k + \beta \mathbf{v}_k)$$

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{v}_{k+1}$$



Adagrad

- Instead of using the same learning rate for all components of \mathbf{x} ,
Adaptive Subgradient method (Adagrad) adapts the learning rate for each component of \mathbf{x} .
 For each component of \mathbf{x} , the update equation is

$$x_{i,k+1} = x_{i,k} - \frac{\alpha}{\epsilon + \sqrt{s_{i,k}}} \nabla f_i(\mathbf{x}_k)$$

where

$$s_{i,k} = \sum_{j=1}^k (\nabla f_i(\mathbf{x}_j))^2$$

$$\epsilon \approx 1 \times 10^{-8}, \alpha = 0.01$$

- components of s are strictly nondecreasing, hence learning rate decreases over time

RMSProp

- Extends Adagrad to avoid monotonically decreasing learning rate by maintaining a decaying average of squared gradients

$$\hat{s}_{k+1} = \gamma \hat{s}_k + (1 - \gamma) (\nabla f(\mathbf{x}_k) \odot \nabla f(\mathbf{x}_k)), \quad \gamma \in [0, 1], \quad \odot \text{ element-wise product}$$

Update Equation

$$\begin{aligned} x_{i,k+1} &= x_{i,k} - \frac{\alpha}{\epsilon + \sqrt{\hat{s}_{i,k}}} \nabla f_i(\mathbf{x}_k) \\ &= x_{i,k} - \frac{\alpha}{\epsilon + RMS(\nabla f_i(\mathbf{x}_k))} \nabla f_i(\mathbf{x}_k) \end{aligned}$$

root mean square: For n values $\{x_1, x_2, \dots, x_n\}$

$$x_{RMS} = \sqrt{\frac{1}{n} (x_1^2 + x_2^2 + \dots + x_n^2)}.$$

AdaDelta

Also extends Adagrad to avoid monotonically decreasing learning rate
Modifies RMSProp to eliminate learning rate parameter entirely

$$x_{i,k+1} = x_{i,k} - \frac{RMS(\Delta x_i)}{\epsilon + RMS(\nabla f_i(\mathbf{x}))} \nabla f_i(\mathbf{x}_k)$$

Adam

- The **adaptive moment estimation method** (Adam), adapts the learning rate to each parameter.
- stores both an exponentially decaying gradient like momentum and an exponentially decaying squared gradient like RMSProp and Adadelta
- At each iteration, a sequence of values are computed

Biased decaying momentum

$$\mathbf{v}_{k+1} = \beta \mathbf{v}_k - \alpha \nabla f(\mathbf{x}_k)$$

Biased decaying squared gradient

$$\mathbf{s}_{k+1} = \gamma \mathbf{s}_k + (1 - \gamma) (\nabla f(\mathbf{x}_k) \odot \nabla f(\mathbf{x}_k))$$

Corrected decaying momentum

$$\hat{\mathbf{v}}_{k+1} = \mathbf{v}_{k+1} / (1 - \gamma_{v,k})$$

Corrected decaying squared gradient

$$\hat{\mathbf{s}}_{k+1} = \mathbf{s}_{k+1} / (1 - \gamma_{s,k})$$

Next iterate

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha \hat{\mathbf{v}}_{k+1} / (\epsilon + \sqrt{\hat{\mathbf{s}}_{k+1}})$$

- Defaults: $\alpha = 0.001$, $\gamma_v = 0.9$, $\gamma_s = 0.999$, $\epsilon = 1 \times 10^{-8}$

Adamax

Same as Adam, but based on the max-norm L_∞ .

$$\begin{aligned}\mathbf{s}_{k+1} &= \gamma^\infty \mathbf{s}_k + (1 - \gamma^\infty) (\|\nabla f(\mathbf{x}_k)\|_\infty) \\ &= \max(\gamma \mathbf{s}_k, \|\nabla f(\mathbf{x}_k)\|_\infty)\end{aligned}$$

Nadam

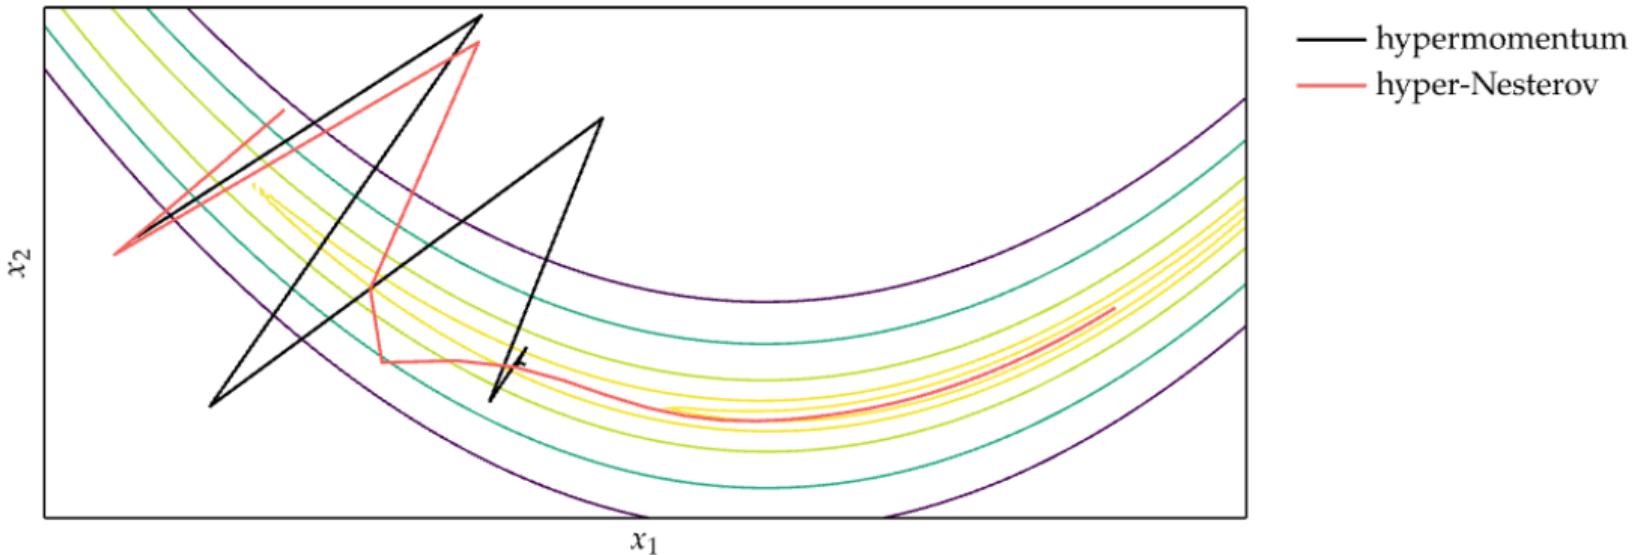
Nadam

- Nesterov-accelerated Adaptive Moment Estimation
- Adam is basically RMSProp with momentum
- We have seen that Nesterov is often more efficient
- Welcome to Nadam: Adam which uses the Nesterov momentum.

Hypergradient Descent

- Learning rate determines how sensitive the method is to the gradient signal.
- Many accelerated descent methods are highly sensitive to hyperparameters such as learning rate.
- Applying gradient descent to a hyperparameter of an underlying descent method is called hypergradient descent
- Requires computing the partial derivative of the objective function with respect to the hyperparameter

Hypergradient Descent



Summary

- Gradient descent follows the direction of steepest descent.
- The conjugate gradient method can automatically adjust to local valleys.
- Descent methods with momentum build up progress in favorable directions.
- A wide variety of accelerated descent methods use special techniques to speed up descent.
- Hypergradient descent applies gradient descent to the learning rate of an underlying descent method.