

Advanced Machine Learning


Neural Networks and Gradient-Based Learning

Lukas Galke
Spring 2026

Keywords for this session


- Neural Networks
- Gradient-based Learning

The Neuron ... in Nature



That is nice ... but


- We will use Neurons as basic simple "prediction machines"
- What does that mean?




Example




Example




Example



Example: Learning



Example: Learning




Now with more inputs

- This was extremely simplistic
- Now let's look at a more complicated thing:
 - You are thinking whether you should attend a sports event
- You base the decision on the following factors:
 1. Temperature?
 2. Ticket Price?
 3. Travel time?


Now with more inputs

- This was extremely simplistic
- Now let's look at a more complicated thing:
 - You are thinking whether you should attend a sports event
- You base the decision on the following factors:
 1. Temperature? x_1
 2. Ticket Price? x_2
 3. Travel time? x_3

Now with more inputs




Now with more inputs



➤ How to make the decision?

$$x_1 + x_2 + x_3 > \text{threshold?}$$

Now with more inputs




➤ How to make the decision?

$$x_1 + x_2 + x_3 > \text{threshold}?$$

➤ That would mean, the higher the price and the farther away, the more likely you are to go ...

Now with more inputs




- How to make the decision?
- Therefore, introduce weights:

$$W_1x_1 + W_2x_2 + W_3x_3 > \text{threshold?}$$



- Now, you can set the weights and threshold according to your preference, e.g., 1, -4, -1

This is in fact how neurons work



- An artificial neuron (aka perceptron) consists of
 - A number of weighted inputs
 - An activation function
 - The generated output


From Neuron to "Brain"



- We now can connect neurons
- The output of one neuron becomes the input of other neurons

In a more structured way

- We normally have more than one node
- Multiple Nodes are arranged in layers
- Each layer receives the generated output from the previous layer



How to build these networks?

➤ Define the Architecture

- How many inputs?
- How many hidden layers?
 - How many neurons per layer?
 - What activation function to use?
- What is the desired output? Define output neurons
 - One-out-of-N classification: Softmax
 - Multiple-out-of-N classification: Sigmoid
 - Regression: linear output

➤ Define the loss function

- Dependent on the output and task

➤ Train the network

How is a Network Trained?

➤ General Procedure:

- We present the network with an example where we know the answer
- We observe the answer of the network and adjust the weights accordingly

➤ Ideal world:

- We simply look at the dataset and could exactly calculate the weights

➤ Reality:

- We define a **cost-function**, the so-called **loss function**
- We iteratively approximate the best setting by trying to minimize the cost functions
- We do this by a process called **gradient descent**


Keywords for this session

- Neural Networks
- Gradient-based Learning



The Central Idea

- Update the model parameters following the steepest slope of the loss function



More Mathematically

- Suppose function $y = f(x)$
- Derivative of function denoted: $f'(x)$ or as dy/dx
 - Derivative $f'(x)$ gives the slope of $f(x)$ at point x
 - It specifies how to scale a small change in input to obtain a corresponding change in the output:
$$f(x + \varepsilon) \approx f(x) + \varepsilon f'(x)$$



- We know that

$$f(x - \varepsilon \text{sign}(f'(x)))$$

is less than $f(x)$ for small ε .

- Thus we can reduce $f(x)$ by moving x in small steps with opposite sign of derivative
- This technique is called gradient descent (Cauchy 1847)

Gradient Descent



The General Case

- Let us assume we have a data dataset $X = \{x_1, \dots, x_N\}$
- We have defined a neural network with the parameters θ
- We further have defined a cost function J :

$$J(\theta, x, y)$$

- Which give us per sample x and the true label y , and the current parameters of the model a certain cost
 - If we, e.g., misclassify x we will get high costs, if we are correct, very low costs

The General Case

- For each object x_1 we can now calculate the gradient:

$$g = \nabla_{\theta} J(\theta, x_1, y_1)$$

- This gradient tells us how to modify the weights in order to achieve a lesser weight, but is noisy and overly depended on x_1
- Since we want to improve on the entire dataset, we calculate the overall gradient:

$$g = \frac{1}{N} \sum_X \nabla_{\theta} J(\theta, x_i, y_i) = \nabla_{\theta} J(\theta, X, y)$$

A bit more formally


➤ We now modify the weights according to the gradient descent method:

$$\mathbf{g} = \frac{1}{N} \sum_{\mathcal{X}} \nabla_{\theta} J(\theta, \mathbf{x}_i, \mathbf{y}_i) = \nabla_{\theta} J(\theta, \mathbf{X}, \mathbf{y})$$

$$\theta_{\text{new}} = \theta_{\text{old}} - \epsilon \mathbf{g}$$

- ϵ is called the learning rate
- Note, most of the time, we do not compute the gradient for all available data, but for randomly selected small portions of the dataset.
- This is called **Stochastic Gradient Descent**

Gradient Descent



The Loss Function

- We are often using the loss function called cross-entropy

$$H(p, q) = - \mathbb{E}[\log(q)]$$

- For most of our (discrete) cases

$$H(p, q) = - \sum_{x \in X} p(x) \log(q(x))$$

- Where $p(x)$ is the real data distribution, while $q(x)$ is our prediction


- Interpretation: The average number of bits needed to encode data distributed according to $p(x)$, if you assume it follows $q(x)$.

Keywords for this session

- Neural Networks
- Gradient-based Learning



Outlook



How's my teaching?



<https://www.admonymous.co/lukasgalke>