

Figure 6. Brazil state map of ChatGPT's ranking of “Where is smarter”.

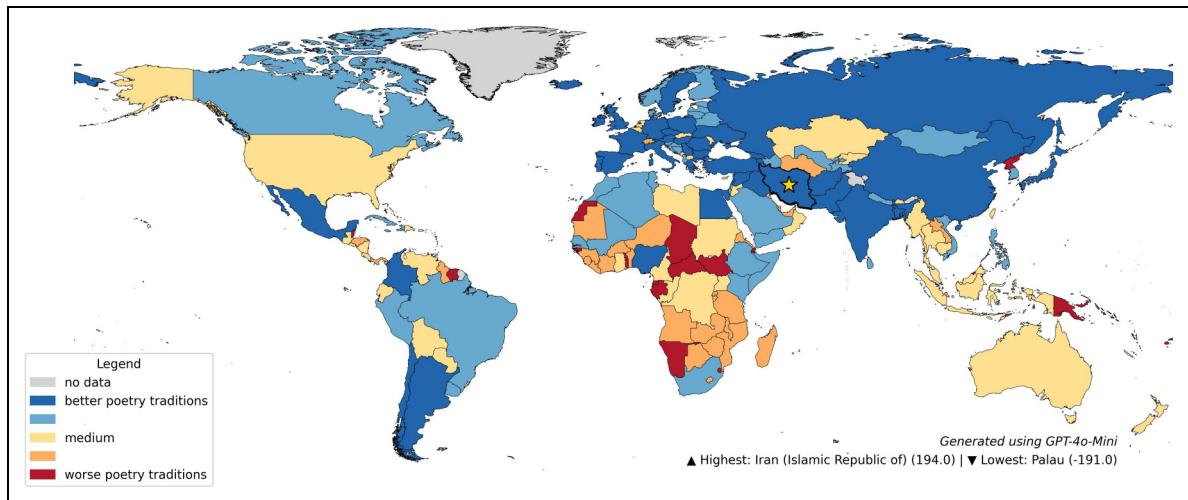


Figure 7. Country map of ChatGPT's ranking of “Where has better poetry traditions”.

not trigger content moderation systems, they persistently echo racialised, gendered, or colonial imaginaries. Because such patterns are widespread and not explicitly hateful, they can pass through filters and become amplified

through repetition. In short, these outputs highlight a representational problem: they echo familiar moral caricatures, granting stereotypes renewed authority through the model's patterned repetition.

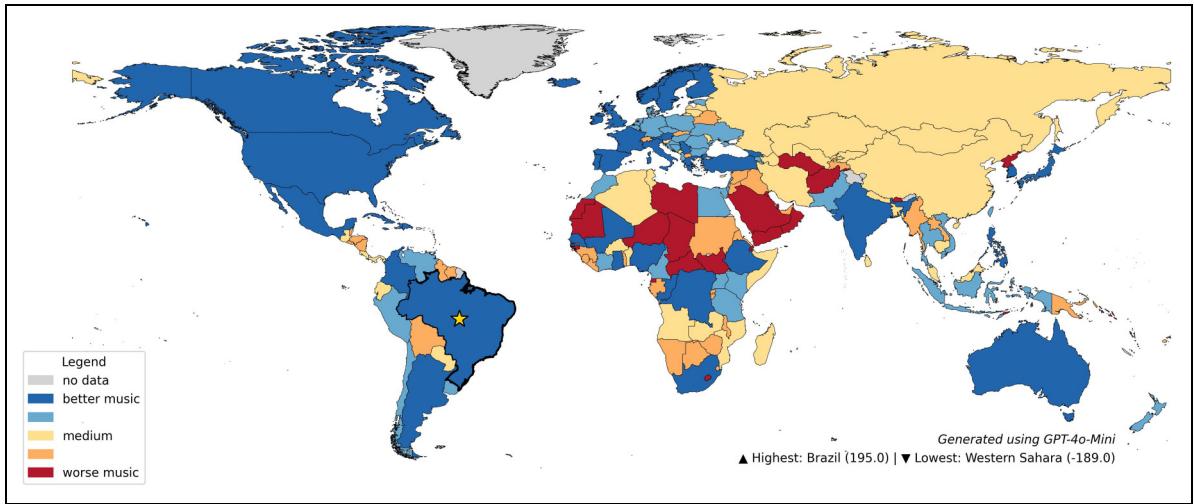


Figure 8. Country map of ChatGPT's ranking of “Where has better music”.

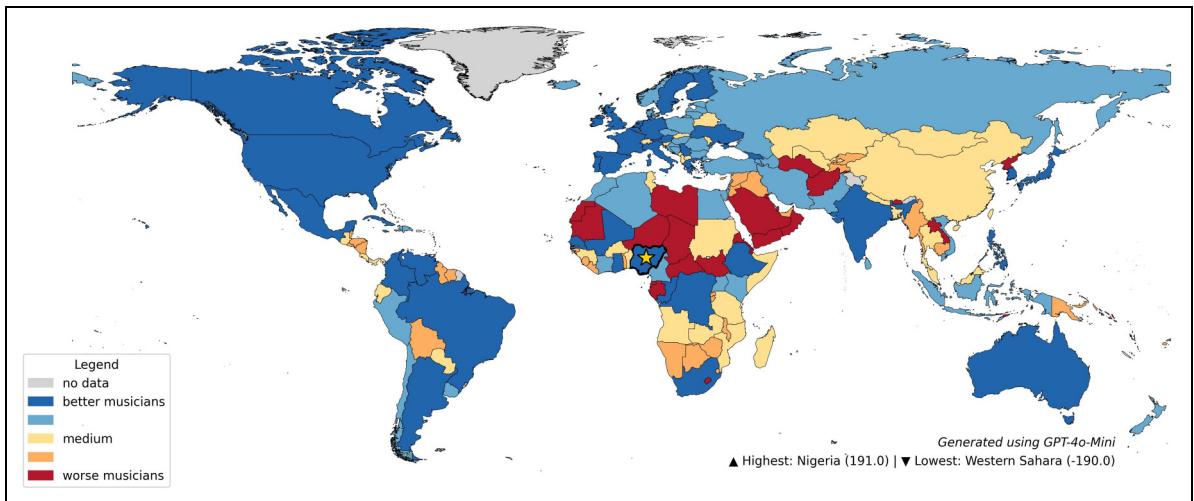


Figure 9. Country map of ChatGPT's ranking of “Where has better musicians”.

“Stinginess” is a well-worn trope (e.g. Scots and Dutch) that resurfaces in language models because it appears frequently but is rarely flagged as hate speech. This is the core of trope bias: shallow cultural characterisations re-emerge as plausible facts through patterned repetition in the training data. Figure 10 shows ChatGPT’s ranking of stinginess, revealing how the model compensates for a lack of standardised data by leaning on stereotypical cues likely drawing on uneven media narratives and inherited caricatures. North Korea, for example, is ranked the stingiest despite limited reliable data, likely a product of its negative coverage in general. In the case of Venezuela (also highly ranked), years of crisis have generated numerous depictions of scarcity that cause the silicon gaze to amplify simplistic tropes to fill informational gaps, re-casting caricatures in the guise of insight.

Trope bias is especially evident in responses to open-ended prompts like “Which country has better vibes?” – a question intentionally selected for its vagueness and cultural subjectivity. Without any formal metric for “vibes,” the model turns to familiar slogans and high-frequency media tropes. Costa Rica tops the list (see Figure 11), almost certainly due to the global circulation of its “pura vida” ethos, which appears frequently in travel writing and social media as shorthand for relaxed, happy living. At the other end of the spectrum, North Korea is ranked as having the worst vibes, not due to much direct discussion of “vibes” in relation to the country, but geopolitical narratives about famine and repression. Here, trope bias seeks easy archetypes: good vibes become Costa Rica’s slogans and bad vibes become authoritarianism and isolation, recycling surface-level associations as seemingly objective judgments.

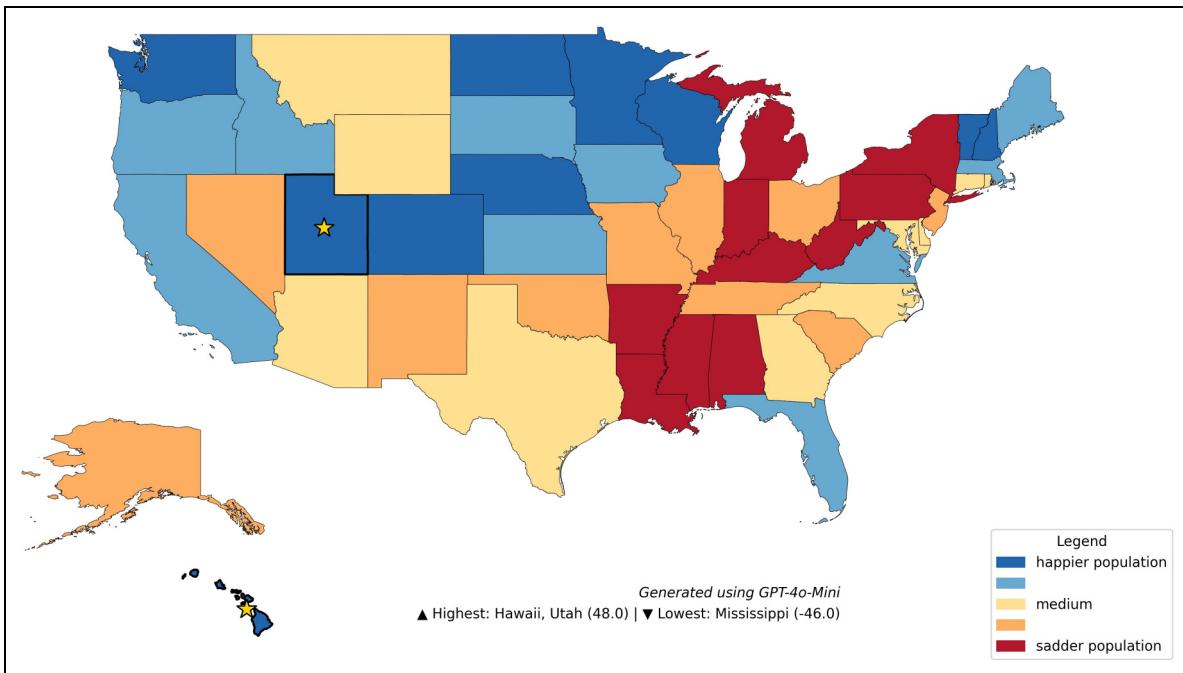


Figure 13. US state map of ChatGPT's ranking of “Where has a happier population”.

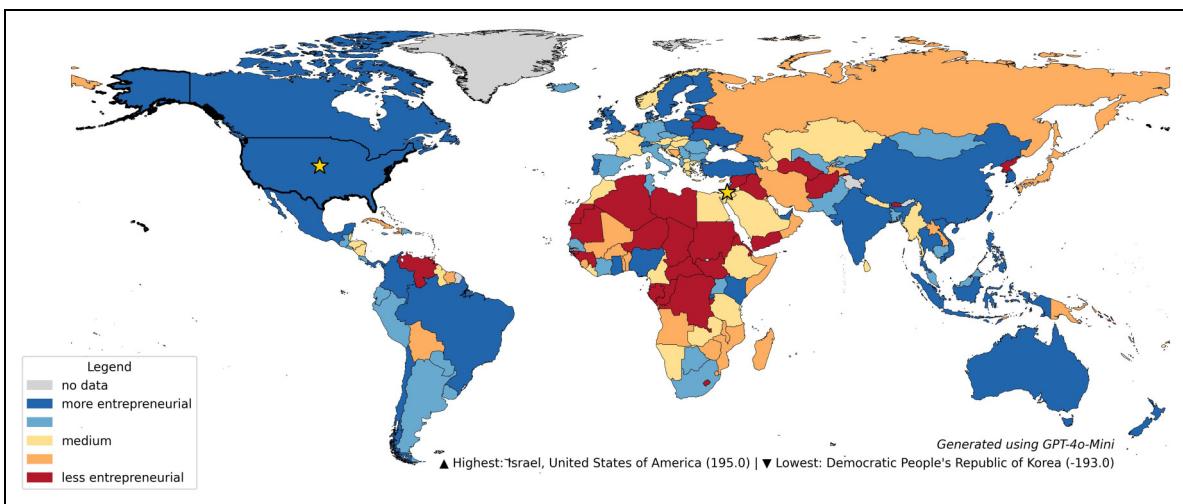


Figure 14. Country map of ChatGPT's ranking of “Where has more entrepreneurial spirit”.

embedded innovation, common in many Global South contexts, escape the metric's gaze and are therefore treated as evidence of entrepreneurial absence. Taken together, these cases show proxy bias to be more than a statistical artefact; it is a political project of legibility (Scott, 1998) in which models privilege what can be counted, certify it as common sense, and deepen existing asymmetries in global visibility. Because these models learn exclusively from text rather than structured or tabular data, the prominence of certain statistics in their outputs reflects how often those figures appear in their source documents. In practice, this means that

numerical data from well-covered places will be much more likely to dominate the model's responses, while data for less-reported countries will be even further underrepresented. In other words, the model's ‘view’ of the world mirrors the uneven frequency with which different regions are discussed in its training texts.

Conclusion

In this article, we document how the silicon gaze exhibits biases across different geographies and how generative AI

